

## Concurrency Control

Jatin

(1)

Two Phase  
locking

Snapshot (Time Stamp Base 2  
Protocol)

## Lock Based Protocol

One way to ensure isolation is to require that data items be accessed in a mutually exclusive manner. The most common method used to implement this requirement is to allow transaction to access a data item only if it is currently holding a lock on that item.

### Locks (Binary Locks)

There are various modes in which data item may be locked.

→ Shared mode lock (S),  $T_i$  can read but cannot write.

→ Exclusive mode lock (X):  $T_i$  can both read & write.

	S	X
S	True	False
X	False	False

Lock Compatibility Matrix

$T_1$ :  
 lock - x(B);  
 R(B);  
 $B := B - 50;$   
 w(B);  
 unlock(B);  
 lock - x(A);  
 R(A);  
 $A := A + 50;$   
 w(A);  
 unlock(A);

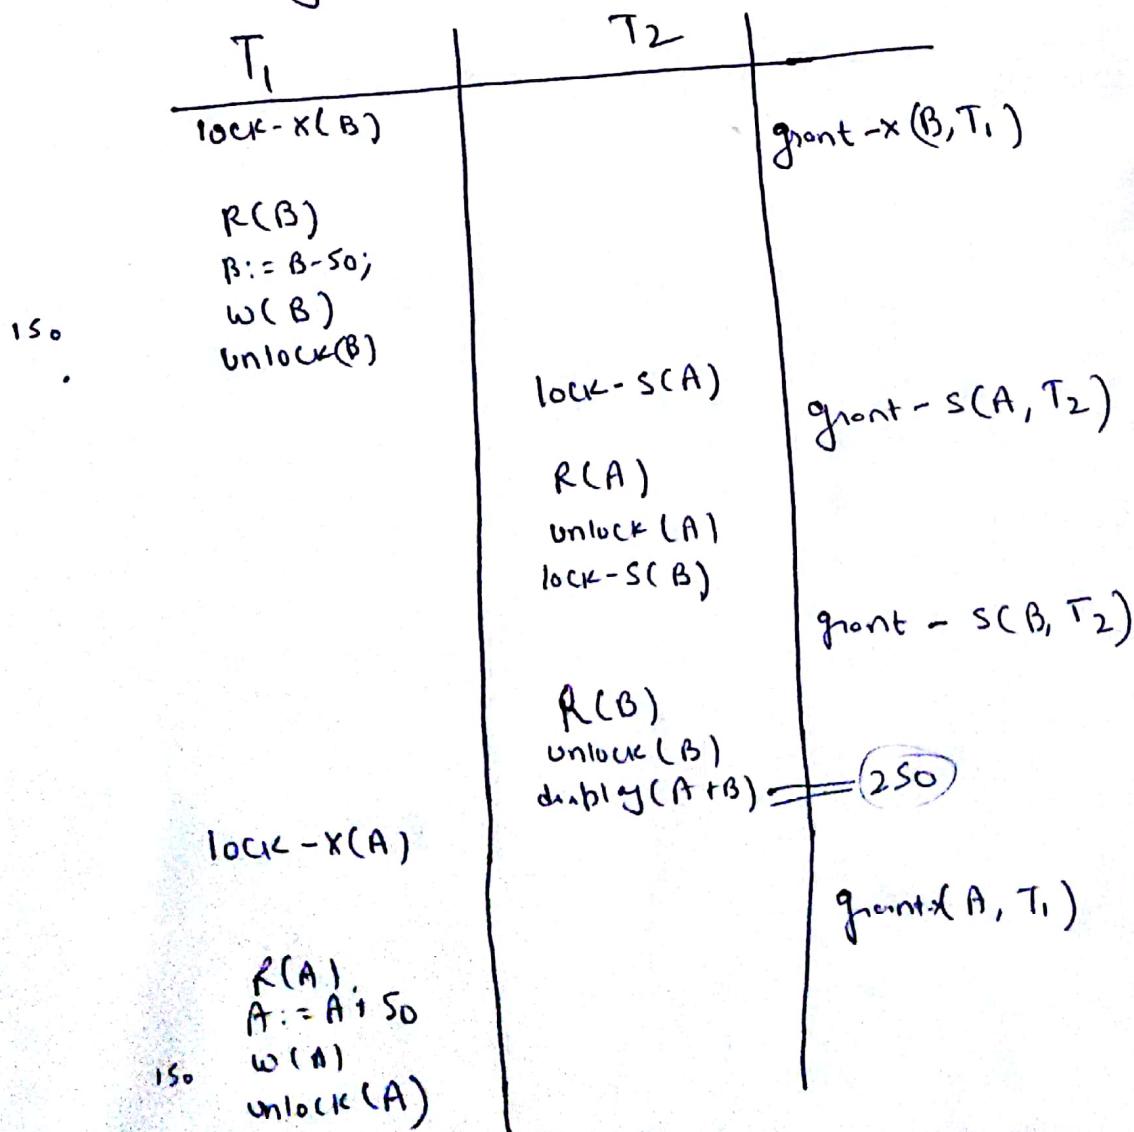
$T_2$ :  
 lock - s(A);  
 R(A);  
 unlock(A);  
 lock - s(C(B));  
 R(B);  
 unlock(B);  
 display(A+B);

display sum of A+B

### Transfer Balance

Now if executed serially  $T_1, T_2$ , or  $T_2, T_1$  the sum displayed is consistent. But if execute concurrently  
the accordig following Schedule 1.

$$\begin{aligned} A &= 100 \\ B &= 200 \end{aligned}$$



Now this Schedule is giving incorrect result  
 of  $A + B$  which occurs due to early release of lock and this operation can be delayed as per the modified Transaction  $T_1'$

```

lock - X(B);
R(B);
B := B - 50;
W(B)
lock - X(A);
R(A);
A := A + 50;
W(A);
unlock(B);
unlock(A);
    
```

### Problem

inconsistency, deadlock  
 Not guaranteed serializability

Now  $T_2$  will print correct result of  $A + B$ ,  
 because unlocking is delayed.

### Deadlock

$T_3$	$T_4$
$\text{lock - } X(B)$ $R(B)$ $B := B - 50$ $W(B)$	$\text{lock - } S(A)$ $R(A)$ $lock - S(B)$ $lock - x(A)$

Consider the partial schedule created by  $T_3, T_4$   
 $T_3$  is requesting Shared lock on B which is held by  $T_3$ .

$T_3$  is requesting Exclusive lock on A which is held by  $T_4$  and causes deadlock. It can overcome by rollback.

It is desirable over inconsistent state of DB. If locking don't use or Unlocking is done earlier, then DB leads to inconsistent state.

## Locking Protocol

It indicates when a transaction may lock and unlock each of the data items. Locking protocol restricts the number of possible schedules.

### Granting of locks

When data item is not locked by any transaction, the lock can be granted. However care must be taken while granting locks. Consider a scenario,  $T_2$  has shared mode lock on B, &  $T_1$  is requesting exclusive mode lock on B.  $T_1$  has to wait. Meanwhile  $T_3$  may request shared lock on B and since it is compatible, therefore it is granted. Now  $T_1$  has to further wait for  $T_3$ . This leads to starvation. It can be prevented by following ways:-

- (1) There is no other transaction holding a lock on B in a mode that conflicts with Mode
- (2) There is no other transaction that is waiting for a lock on data item and that mode is lock request before  $T_1$ .

## Two Phase Locking Protocol

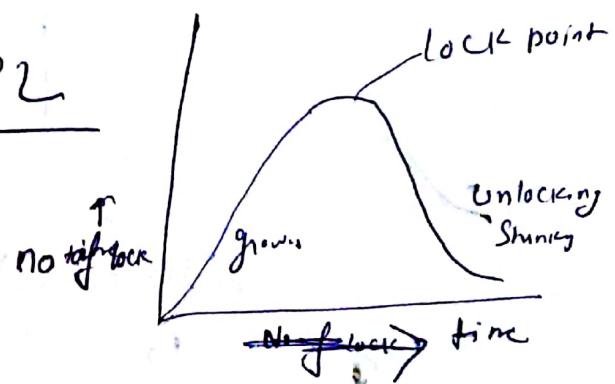
+ ensures serializability by requiring that each transaction issue lock and unlock request in two phase.

Growing phase : Transaction may obtain locks, but may not release any lock.

Shrinking phase : A transaction may release locks, but may not obtain any new locks.

$T_1$  :  
 lock -  $x(B)$ ;  
 $R(B)$ ;  
 $B := B - 50$ ;  
 $w(B)$ ;  
lock -  $x(A)$ ;  
 $R(A)$ ;  
 $A := A + 50$ ;  
 $w(A)$ ;  
unlock ( $B$ );  
unlock ( $A$ );

2 PL



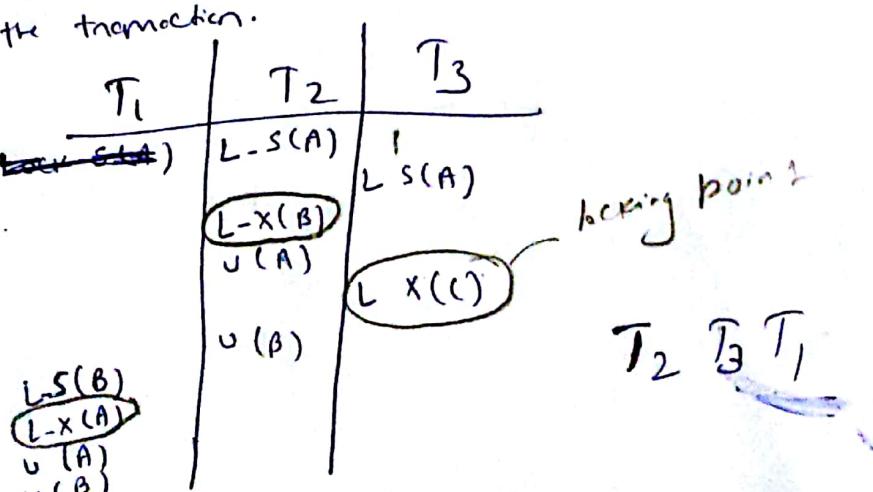
The point in the schedule when the transaction has obtained its final lock (the end of its growing phase) is called the lock point of the transaction. Now if transactions are ordered according to their lock points, then it ensures

serializability for the transaction.

$T_1$	$T_2$
<del>lock</del>	
$l_x(A)$	$l_s(A)$
$w(A)$	
$w(B)$	

$T_1$	$T_2$	$T_3$
<del>lock</del>	$l_s(A)$	$l_s(A)$
	$l_x(B)$	$l_x(C)$
	$u(A)$	
	$u(B)$	



→ Two phase locking does not ensure freedom from deadlock.

	$T_3$	$T_4$
	$L - X(B)$	
	$R(B)$	
	$B = B - S_0$	
	$W(B)$	
		$Lock - S(A)$
		$R(A)$
		$L - S(B)$
	$L - X(A)$	

locking  
phase

dead  
lock

time

lock

<p style="position: absolute; right: 10%; top: 4

## 2 PL

Reading Rollback can be avoided by a modification of 2PL called Strict two phase Locking.

This protocol follows 2PL as well

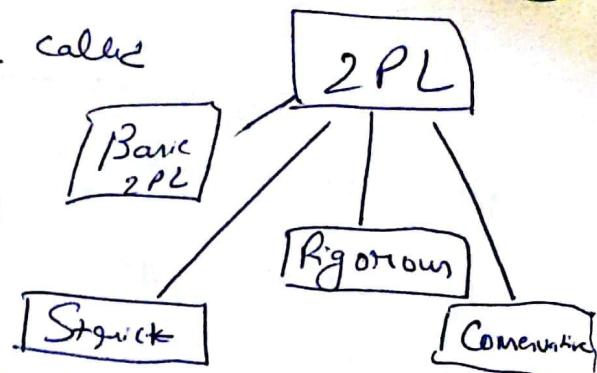
as all exclusive-mode locks

taken by a transaction be held until that

transaction commits, preventing any other transaction

reading that data.

deadlock may occur, And 2  
conceding rollback.



## Rigorous 2PL

It requires that all locks be held until the transaction commits. Transaction can be serialized in the order in which they commit. And Conceding rollback, deadlock may occur

## Conservative

It requires a transaction to lock all the items it acquires before the transaction begins execution, by predeclaring its Read-Set and Write-Set. Read-Set is the set of all items that the transaction reads and similarly write-set is the set of items that ~~all the~~ it writes. If any of the predeclared item is not available for locking, then transaction does not hold any lock, instead it waits till all items are available. It is not mortally possible. Resource utilization is less.

No Conceding rollback, deadlock free

## Lock conversions

It is refinement over basic 2PL, in which lock conversions are allowed. May be upgraded by conversion from shared to exclusive, downgraded by conversion from exclusive to shared. Upgrading allowed only in growing phase and downgrading allowed only in shrinking phase.

$T_8$	$T_9$
$L-S(a_1)$	
$L-S(a_2)$	$L-S(a_1)$
$L-O-S(a_3)$	$Io-S(a_2)$
$L-S(a_4)$	
$L-S(a_n)$	$UnL(a_1)$
$Upgrade(a_1)$	$UnL(a_2)$

It also ensures conflict serializability schedules, transaction can be made serializable by these lock points.

- All the locks held by trans are released after commit / abort

→  $R(Q)$  issues  $L-S(Q)$  followed by  $R(Q)$

→  $w(Q)$  issues  $L-X(Q)$  followed by  $w(Q)$ ; if lock already taken by Trans on  $Q$  as Shared lock then  $Upgrade(Q)$  is used.

## Deadlock

(5)

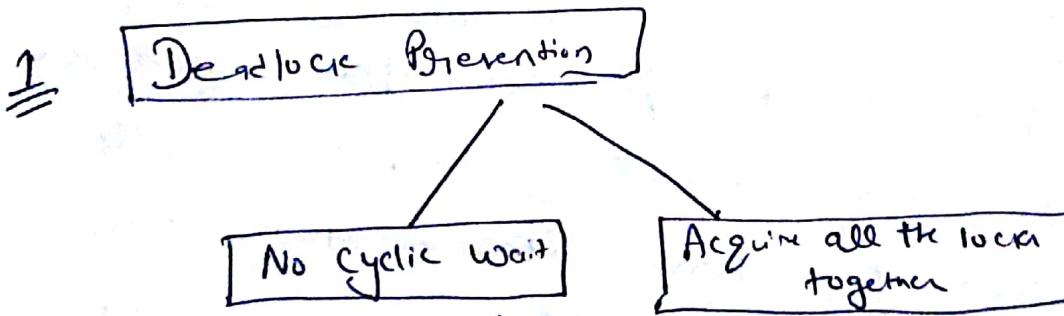
A System is in deadlock state if every transaction in the system is waiting for another transaction in the Set.

$(T_0, T_1, \dots, T_n)$   $T_0$  is waiting for data item held by  $T_1$ ,  $T_1$  is waiting for  $T_2$  and  $T_n$  is waiting for  $T_0$ .

One solution is to rollback some transaction fully or Partial.

Two methods for dealing with deadlock are

- 1 Deadlock Prevention
- 2 Deadlock Detection / Deadlock Recovery



- we total ordering of data items along with 2PL
- Lock must be acquired in the right order
- It is hard to predict, what item to be locked before executing
- Utilization is very low

### Premption and Rollbacks

If  $T_j$  grants the resource held by  $T_i$ , the lock granted to  $T_j$  by preemption and rolling back of  $T_i$ . To control preemptions, a unique timestamp is assigned to each transaction when it begins. This timestamp is used to decide whether a transaction should wait or rollback.

Two deadlock prevention schemes using timestamp + been proposed

### (i) Wait-die

It is non-preemptive technique. When  $T_i$  requests a data item currently held by  $T_j$ ,  $T_i$  is allowed to wait only if it has a timestamp smaller than that of  $T_j$ , otherwise  $T_i$  is rolled back.

for eg  $T_1, T_2, T_3$       if  $T_1$  request data item held by  
TS:-    5    10    15       $T_2$ , then  $T_1$  will wait  
      old                  new      if  $T_3$  request data item held by  
                                     $T_2$ , then  $T_3$  will be rolled back.

### (ii) Wound-wait

It is preemptive technique, counterpart of wait-die. When  $T_i$  request data item held by  $T_j$ ,  $T_i$  is allowed to wait only if it has a timestamp larger than that of  $T_j$  ( $T_i$  is young than  $T_j$ ), otherwise  $T_j$  is rolled back.

for eg  $\overbrace{T_1 \quad T_2} \quad T_3$       if  $T_1$  request data item held  
TS:-    5:    10    15      by  $T_2$  then data item will be  
                                    preempted from  $T_2$  &  $T_2$  will be  
                                    rolled back. [if  $T_3$  request  
                                    from  $T_2$  then,  $T_3$  will wait.]

(6)

Items with these schemes

→ unnecessary rollbacks

### Lock Timeout Technique

In this scheme, a transaction that has requested a lock waits for at most a specified amount of time. If the lock has not been granted within that time, the transaction is said to be time out, and rolls itself back and restarts. If there was deadlock, one or more transaction stucked in deadlock will time out & roll back and allowing the others to proceed.

- Easy to implement
- Work well with short transactions
- Hard to decide how long a transaction must wait before timeout.
- Long wait is unnecessary once a deadlock occurs, & short wait generally rolls back even when no deadlock.
- Resource wastage
- Starvation.

## Deadlock Detection & Recovery

If prevention is not used, then detection & recovery scheme must be used. It is invoked periodically to check for deadlock. If deadlock occurs, then system must attempt to recover from deadlock.

- Maintain information about current allocation, outstanding requests.
- Use detection technique that uses this information
- Recover from the deadlock.

## Detection Algorithm

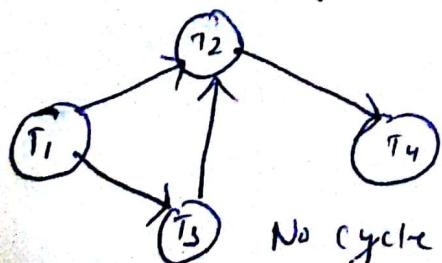
### Wait-for-Graph

$$G = (V, E)$$

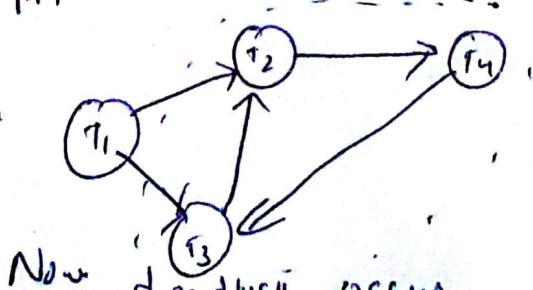
$V$ : Set of Transaction

$E$ : Set of ordered pair  $T_i \rightarrow T_j$ :  $T_i$  is waiting for  $T_j$  to release data item held by  $T_j$ .  
Edge is removed when data item is released by  $T_j$ .

If cycle exists, then deadlock occurs. System maintains WFG, and periodically invoke an algorithm that searches for a cycle in the graph.



Now  $T_2$  holds  
 $T_3$



Now  $T_3$  holds  $T_1$

## Challenges of Detection Algo

- How often does a deadlock occur
- How many transaction will be affected by deadlock

## Recovery from Deadlock

### 1. Selection of Victim

We must determine which transaction to roll back to break the deadlock. Rollback only those transaction which incur minimum cost.

- How long transaction has computed
- How many data items has used
- How many data item needed
- How many transaction will be involved in the rollback.

### 2. Rollback

→ One victim is decided. Simple solution is total rollback. Abort the transaction and then restart it.

Partial Rollback: → Undo till the lock causing deadlock.

3. Starvation: It may happen that same transaction becomes victim based on cost factor. So transaction must be picked up as victim only a finite number of times.

## Time Stamp Based Protocols (Deadlock Preventive Protocol) ⑧

Earlier scheme utilizes scheme of ordering transaction based every conflicting pair of transactions at execution time. This technique utilizes an ordering among transactions based on timestamp-ordering.

Associate each transaction  $T_i$  with a unique fixed timestamp denoted by  $TS(T_i)$ :

If  $TS(T_i) < TS(T_j)$  then  $T_j$  is new.

Timestamp may be implemented by two methods

↳ System Clock: Timestamp value is equal to value of clock

↳ Logical Counter: A counter value is incremented after a new timestamp has been assigned.

The timestamps of transaction determines the serializability order. Thus if,  $TS(T_i) < TS(T_j)$  then the system must ensure that the produced schedule is equivalent to a serial schedule.

Assign Two TS values with each data item Q:

w-timestamp( $\alpha$ ): largest timestamp of any transaction that executed write( $\alpha$ ) successfully

R-timestamp( $\alpha$ ): largest timestamp of any transaction that executed Read( $\alpha$ ) successfully.

## Timestamp Ordering Protocol

It ensures that any conflicting read and write operations are executed in timestamp order. It follows the following operation.

### $T_i$ issues Read( $Q$ )

- If  $TS(T_i) < w\text{-timestamp}(Q)$  then  $T_i$  needs to read a value of  $Q$  that was already overwritten - then Read operation is rejected and  $T_i$  is rolled back.
- If  $TS(T_i) \geq w\text{-timestamp}$ , then read is executed and  $R\text{-timestamp}$  is set to maximum of  $R\text{-timestamp}(Q)$  and  $TS(T_i)$ .

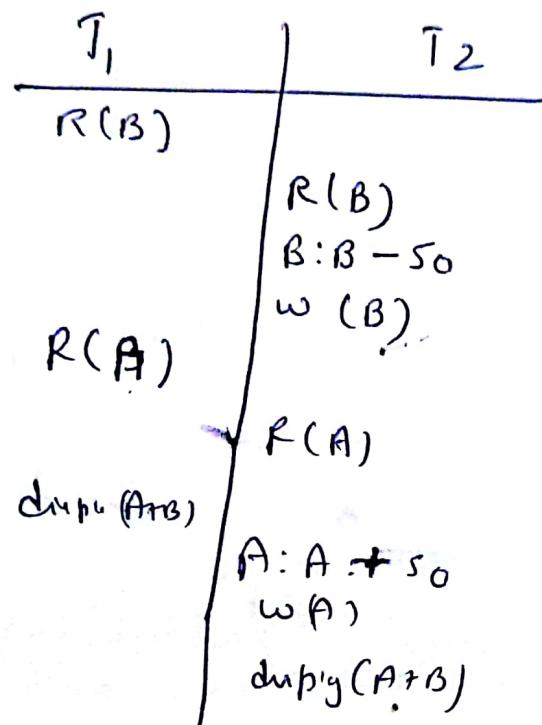
### $T_i$ issues Write( $Q$ )

- If  $TS(T_i) < R\text{-timestamp}(Q)$ , then the value of  $Q$  that  $T_i$  is producing was needed previously and the system assumed that value would never be produced. Hence system rejects the write operation & rollback.
- If  $TS(T_i) < w\text{-timestamp}(Q)$  then  $T_i$  is attempting to write an obsolete value of  $Q$ , hence system rejects write operation and rollback.
- Otherwise, System executes the write operation & sets  $w\text{-timestamp}(Q)$  to  $TS(T_i)$

(9)

$T_1$   
 $R(B)$   
 $R(A)$   
 $\text{Dupl}(A+B)$

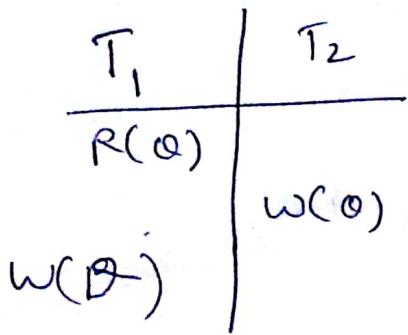
$T_2$      $R(B);$   
 $B := B - 50;$   
 $W(B);$   
 $R(A);$   
 $A := A + 50;$   
 $W(A)$   
 $\text{Dupl}(A+B);$



This protocol can generate schedules that are not recoverable. However it can be extended to make the schedules recoverable in one of several ways.

- Recovery / cascadeliveness can be ensured by performing all writes together at the end of the transaction.
- Recovery / cascadeliveness can also be guaranteed by using a limited form of locking, whereby reads of uncommitted items are postponed until the transaction that updated the item commits.
-

## Thomas Write Rule



$$TS(T_1) < TS(T_2)$$

$R(Q)$  of  $T_1$  succeed and  $w(Q)$  of  $T_2$  succeeded.

When  $w(Q)$  of  $T_1$  attempts to operate, then

$TS(T_1) < w\text{-timestamp}(Q)$  Thus it is rejected and  $T_1$  is rolled back.

Although  $T_2$  has already wrote  $Q$ , and roll back of  $T_1$  is unnecessary and the value  $T_1$  is going to write in  $Q$  is never read. So any transaction  $T_i$  whose  $TS(T_i) < TS(T_2)$  that attempts to read ( $Q$ ) will be rolled back.

And any transaction  $T_j$  whose  $TS(T_j) > TS(T_2)$  must read the value of  $Q$ .

So Thomas' write operation can be ignored under certain circumstances - Only 2nd rule of write is modified.

If  $TS(T_i) < w\text{-timestamp}(Q)$  then  $T_i$  is attempting to write on obsolete value of  $Q$ . Hence, this write operation can be ignored.